



Enabling Reproducibility with Docker

Enabling Reproducibility with Docker

- Daniel Trahan
 - *Email:* Daniel.Trahan@Colorado.edu
 - *RC Homepage:* <https://www.colorado.edu/rc>
 - *RC Email:* rc-help@colorado.edu
-
- Slides available for download at:
https://github.com/ResearchComputing/Containers_Fall_2021

Outline

- Part 1: Container fundamentals and Docker (9/22/21)
 - Reproducibility and the Case for Containers
 - Containers
 - Docker
 - Images and Containers
 - Commands
 - File Access
 - Building Docker Images
 - Dockerhub
- Part 2: Containers for HPC w/ Singularity (9/29/21)

Demo Files:

- Before we begin I will note that this tutorial will have interactive components.
- If you would like to participate in the demos provided for this tutorial then go ahead and clone the test files from the Github.
 1. Navigate to a desired directory
 2. Clone the repository
`git clone https://github.com/ResearchComputing/Containers_Fall_2021`
 3. Navagate into the directory and store the path into a variable.
`cd Containers_Fall_2021`
`export CONTAINER_ROOT=$(pwd)`

Reproducibility and Research

- Scientific Software is often challenging to work with
 - Difficult installation
 - Low support from the developers
 - Very outdated
 - Complex Dependency trees
- Because of this its often desired for a software to be repeatable and accurate.
- *But installs are only done once. Why should I care about reproducible applications.*

Reproducibility and Research

- Research is Collaborative
 - Team members work together to get projects done.
 - Reproducibility ensures all members of a team can provide productivity towards a project.
- Research is Correcting
 - Research is hard
 - Academic reviews are commonplace
 - Someone may wish to accurately reproduce your work
- Research is Continuous
 - You may be working on a single project for a long period of time
 - What happens in you move, but bring your work to another system?

Options for reproducibility

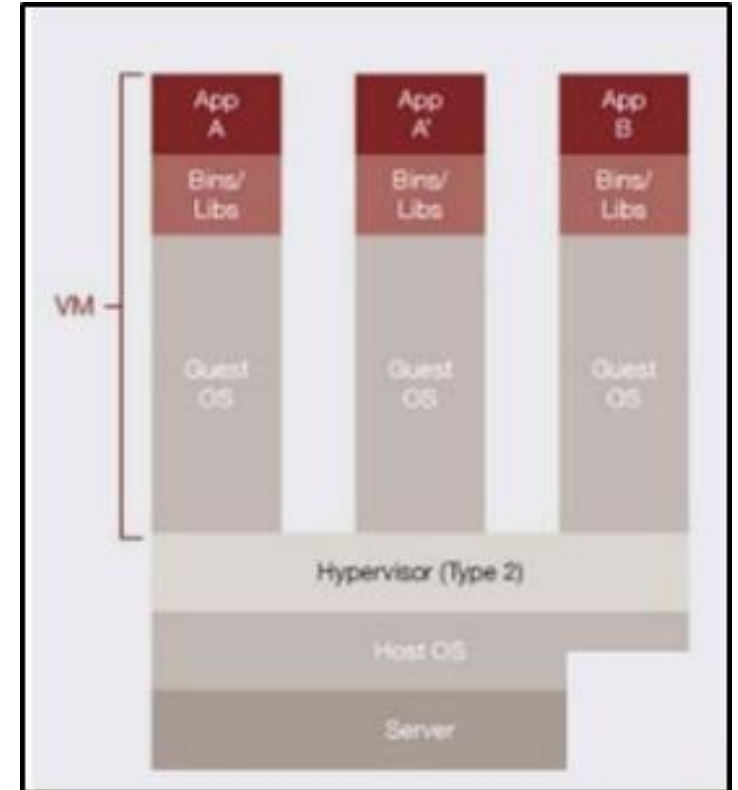
- Lots of options!
 - Detailed instructions
 - Software bundles
 - Virtual Environments
 - Python, Anaconda, Spack
- But do they really enable accurate reproducibility?
 - Incorrect installs?
 - Hardware or OS?
 - Performance?

Containers

- A Container is a packaged bundle of OS, libraries, software and files that runs as a process under a host OS
- Containers use an application on the host operating system called a Container Manager
 - Manages operating system and libraries run as containers
 - Like virtual machines, but does not need dedicated CPUs memory or storage

Virtualization (1)

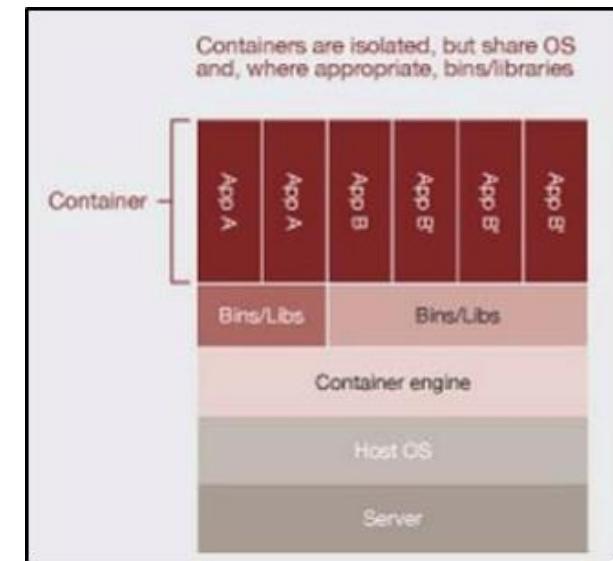
- Virtualization is a technology that utilizes software to abstract components of a technology
- The most common application is in Hardware Virtualization
 - Virtual Machines
 - Partitions off Memory, CPU, GPU, and Storage
 - Runs a virtual OS
 - Runs software on the virtualized machine
 - Examples: VMware, Virtualbox



Material courtesy: M. Cuma, U. Utah

Virtualization (2)

- Another use of virtualization is in OS Level Virtualization
 - Can run many isolated guest OS instances under a host OS
 - This virtualization is what is used by Docker and other container software.
 - Best of both worlds!
 - Isolated environments
 - No hardware partitioning



Material courtesy: M. Cuma, U. Utah

Containerization Software

- Docker
 - Well established – largest user base
 - Has Docker Hub for container sharing
 - Problematic with HPC (Fix incoming!)
- Singularity
 - Designed for HPC
 - Second largest userbase
 - Developed for scientific use
- Charliecloud; Shifter
 - Designed for HPC
 - Based on Docker
 - Less user-friendly

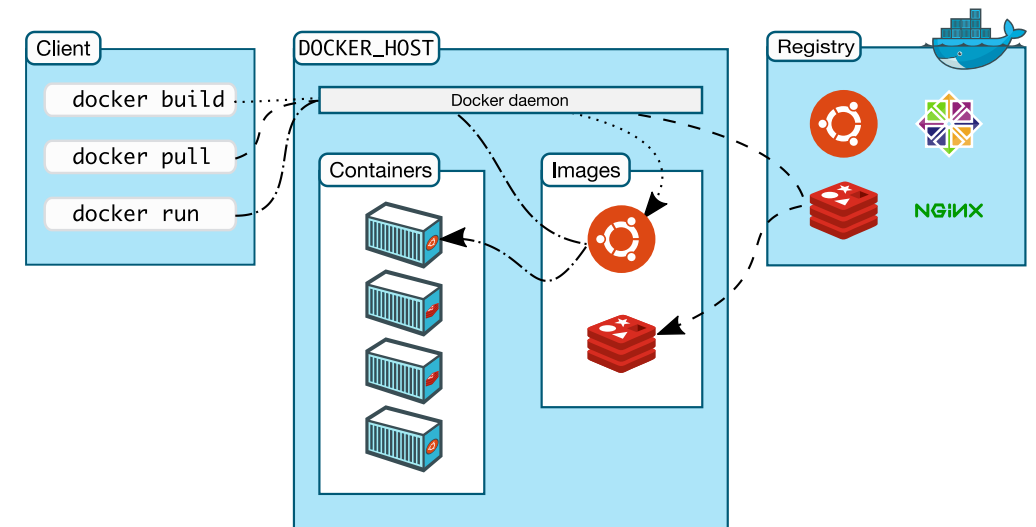


Installing Docker

- Docker Community Edition
 - Comfy GUI to help keep track of containers and images!
 - Available on all operating systems
 - Windows users can enable WSL2 support following the instructions here:
<https://docs.docker.com/docker-for-windows/install/>
- Docker toolbox
 - Legacy solution for Windows and Mac for versions that do not meet the version requirements.
 - Utilizes the Virtual Box hypervisor for virtualization

Docker Nuts and Bolts

- Docker runs on a concept of images and containers.
 - Images: Saved snapshots of a container environment.
 - Made from a Dockerfile or pulled from Docker Hub
 - Stored in the Docker cache on your disk
 - Immutable (mostly...)
 - Containers: Instances of images that are generated by Docker when an image is 'run'
 - Instance of image running in memory
 - Ephemeral and state cannot be saved
 - Can be run interactively



Docker 'Hello World'

- Let's start with something simple:
 - Docker "Hello, World!"
 - Relatively small image
 - No dependencies
 - Built as a general test case
- Command we will run:

`docker run hello-world`

Docker Commands

- Docker Commands are usually in the form of:
`docker <sub-command> <flags> <target/command>`
- Examples:
`docker run -it myimage`
`docker container ls`
`docker image prune`

Launching a Docker Container

- Launch docker image as a container:

```
docker run <image-name>
```

- Run a docker image interactively:

```
docker run -it <image-name>
```

- If an image is not on the system, then Docker will search Dockerhub to see if the image exists.
- Specify commands after your image to execute specific software in your container.

```
docker run <image-name> <program>
```

Exploring a Docker Container

- Docker containers are running tiny operating systems!
- We can explore the operating system by invoking a shell
`docker run -it ubuntu bash`
- This command **launches the ubuntu Docker container with the command 'bash'**

Demo 1: Running a Container

Demo 1: GROMACS

- GROMACS is a molecular dynamics application that can often be a complex and challenging installation for the average user. Linux and Mac only
 - Dense Documentation
 - Software requires compilation
- Luckily, this can be trivialized with Docker!
 - Run the command:
`docker run gromacs/gromacs gmx help commands`
`docker run -it gromacs/gromacs`

Docker Image/Container Commands

| Container Commands | |
|--|---|
| <code>docker container ls</code> | List docker containers currently running: |
| <code>docker container rm <container></code> <code>docker rm <container></code> | Remove (an) container(s): |
| <code>docker container prune</code> | Remove all stopped containers |

| Image Commands | |
|---|-------------------------------------|
| <code>docker image ls</code> | List docker images stored in cache: |
| <code>docker image rm <image></code> <code>docker rmi <image></code> | Remove (an) image(s): |
| <code>docker image prune</code> | Remove unused images |

Docker Image/Container Commands

| Commands | |
|---|---|
| <code>docker info</code> | Shows Docker system-wide information |
| <code>docker inspect <docker-object></code> | Shows low-level information about an object |
| <code>docker config <sub-command></code> | Manage docker configurations |
| <code>docker stats <container></code> | Shows container resource usage |
| <code>docker top <container></code> | Shows running processes of a container |
| <code>docker version</code> | Shows docker version information |

- More details and commands can be found [on the docker documentation page](#)

Dockerhub

- The place where containers live!
- Dockerhub is a Docker hosted library of public and private Docker images.
 - Free and unlimited public images
 - 1 free private repository
- Great for hosting images for fellow researchers
- Commands like git

Building a Docker Container

- To build a docker container, we need a set of instructions Docker can use to set up the environment.
 - Dockerfile
- Once we set up our Dockerfile we can use the command
`docker build -t <image-name> .`
- Then we can run the image with our docker run command
`docker run <image-name>`

Whats in a Dockerfile

- A Dockerfile is simply a text file that contains instructions to build and setup a default Image
 - Commands to build
 - Setting commands
- Requires a source Image

```
mtrahan41@MTrahanRazor15:[ ubuntu-gcc ]$ cat Dockerfile
FROM ubuntu:18.04

RUN apt-get update; \
    apt-get install nano -y; \
    apt-get install gcc -y; \
    mkdir target;

WORKDIR /target
```

Demo 2: Building a Docker Image

Demo 2: Ubuntu w/ GCC

- For this first example we will be building a custom Ubuntu image that will provide a location to run the GNU Compiler Collection.
- Dockerfile provided:
- Need to build:
 1. Navigate to the directory:
`cd $CONTAINER_ROOT/Containers_Fall_2021/dockerdemo/ubuntu-gcc`
 2. Build the image with:
`docker build -t test-gcc .`
 3. Run the image as a container:
`docker run -it test-gcc`

Mounting and File Access (1)

- So now that we have a working container, how can we access the test files we downloaded?
 - Mounting directories: Bind Mount
 - Allows the docker container to access files on the host OS
 - Choose host's *source directory*, files in the directory will be moved to the container's *target directory*
 - **Source Directory**: Directory on the host system. Never within a container.
 - **Target Directory**: Directory in the Docker Container. Never on the host system.
 - A flag set within the docker run command:
`docker run -v <source-dir>:<target-dir> <image>`

Mounting and File Access (2)

- Mounting directories: Volume Mount
 - Same concept, but volumes are stored within docker cache.
 - Create Docker volumes in your terminal and link your volume directory
 - Similarly linked through the docker run command.

```
docker run -v <volume-name>:<target-dir> <image>
```

Demo 2 (Cont.): Mounting

- Returning to our demo, can we give our container access to our test files?
- Let's use a bind mount!
- In the directory where our Dockerfile lives, use this command (all on one line):

```
docker run -it -v $(pwd)/source:/target happy-gcc
```

- We can cd /target and run our test files!
- Command

```
gcc hello.c -o hello.exe  
./hello.exe
```

Modifying a Docker Image

- Suppose you have an existing docker image and want to make changes...
 - Rebuild Dockerfile!
 - Usually a bit cumbersome
- No Dockerfile?
 - Use docker commit!
`docker run -it <image-name> bash # or any shell...`
 - Then commit it to the image
`docker commit <image-id> <image>`

Dockerhub Commands

- Download and upload docker images with ease.

```
docker run <image>
```

```
docker pull <image>
```

- Uploading a little more complicated...

- Sign in with:

```
docker login
```

- List docker images with:

```
docker image ls
```

- Tag your image:

```
docker tag <image-id> <your-username>/<image-name>:<tag>
```

- Push!

```
docker push <your-username>/<image-name>
```

Demo 3: NCL container

Demo 3: NCL Container

- For this next example we will be building a Docker image that will run the NCAR Command Language (NCL)
- Dockerfile provided
- Same process:
 1. Navigate to the Dockerfile found at:
[\\$CONTAINER_ROOT/Containers_Fall_2021/dockerdemo/ncl](#)
 2. Build the Dockerfile and name the image: "ncl-demo"
 3. Run "ncl-demo"
- Can we test a sample script?

Docker Compose

- External Utility that can create and install docker images.
- Builds docker images based on a docker-compose.yml file.
- YAML: YAML Ain't Markup Language
 - Data serialization language
- Describes containers you wish to build with what features.
- Not a docker command but comes bundled with docker!

Docker Compose Commands

- Build all containers in YAML file
`docker-compose build`
- Build and run all containers in YAML file:
`docker-compose up`
- List all containers in YAML file:
`docker-compose images`
- Run a one-off command from a container:
`docker-compose run <container-name> <command>`

Questions?

Additional Resources

- Docker: <https://www.docker.com/>
- Docker Docs: <https://docs.docker.com/>
- Docker Hub: <https://hub.docker.com/>

Thank you!

- Please fill out the survey: <http://tinyurl.com/curc-survey18>
- Contact information: rc-help@Colorado.edu
- Slides: https://github.com/ResearchComputing/Containers_Fall_2021